

Compact Model Deployment: Cost Review

Isacademy Operations, 2026

Summary

Serving costs fell 28 percent quarter over quarter after the compact models replaced the previous 7-billion-parameter endpoint. Average request latency fell from 240 ms to 158 ms. Two incidents were recorded, both caused by cold-start behaviour rather than by model quality.

Contradiction with the research report

Operations measured a 28 percent latency reduction in production, whereas the research report records 34 percent on the validation harness. The gap is attributed to network overhead that the offline harness does not include.

Risks

Cold starts remain the dominant tail-latency risk. Capacity headroom is currently 40 percent and is expected to fall below 20 percent next quarter.